

Detailed Report on MapReduce Implementation

Mohamad Saeed sedighi - 810100179
Mohamad Hosein Motaei - 810199493
Mahdi

April 18, 2025

1 Introduction

The report is structured to include a detailed examination of each file's code, an explanation of its functionality, and an analysis of the test output from running the `Test` script, as depicted in the attached image.

2 Code Analysis

2.1 coordinator.go

```
1 package mr
2
3 import (
4     "log"
5     "net"
6     "net/http"
7     "net/rpc"
8     "os"
9     "sync"
10    "time"
11 )
12
13 type Coordinator struct {
14     mu          sync.Mutex
15     mapTasks    []Task
16     reduceTasks []Task
17     mapDone     bool
18     reduceDone  bool
```

```

19     nReduce      int
20     nMap         int
21     tasks        map[int]*TaskTracker
22 }
23
24 type TaskTracker struct {
25     Status      TaskState
26     WorkerId    string
27     StartTime   time.Time
28 }
29
30 func (c *Coordinator) server() {
31     rpc.Register(c)
32     rpc.HandleHTTP()
33     sockname := coordinatorSock()
34     os.Remove(sockname)
35     l, e := net.Listen("unix", sockname)
36     if e != nil {
37         log.Fatal("listen error:", e)
38     }
39     go http.Serve(l, nil)
40 }
41
42 func (c *Coordinator) monitor() {
43     for {
44         c.mu.Lock()
45         if c.Done() {
46             c.mu.Unlock()
47             return
48         }
49
50         now := time.Now()
51         for _, tracker := range c.tasks {
52             if tracker.Status == InProgress &&
now.Sub(tracker.StartTime) > 10*time.Second {
53                 tracker.Status = Pending
54                 tracker.WorkerId = ""
55                 tracker.StartTime = time.Time{}
56             }
57         }
58         c.mu.Unlock()
59         time.Sleep(time.Second)
60     }
61 }
62
63 func (c *Coordinator) Done() bool {
64     return c.mapDone && c.reduceDone
65 }
66

```

```

67 func MakeCoordinator(files []string, nReduce int)
   *Coordinator {
68     c := Coordinator{
69         mapTasks:    make([]Task, len(files)),
70         reduceTasks: make([]Task, nReduce),
71         tasks:       map[int]*TaskTracker{},
72         nReduce:     nReduce,
73         nMap:        len(files),
74     }
75
76     for i, file := range files {
77         c.mapTasks[i] = Task{Map, i, file, nReduce,
len(files)}
78         c.tasks[i] = &TaskTracker{Pending, "", time.Time{}}
79     }
80
81     for i := 0; i < nReduce; i++ {
82         c.reduceTasks[i] = Task{Reduce, i, "", nReduce,
len(files)}
83         c.tasks[len(files)+i] = &TaskTracker{Pending, "",
time.Time{}}
84     }
85
86     go c.monitor()
87     c.server()
88     return &c
89 }
90
91 func (c *Coordinator) AssignTask(args *TaskArgs, reply
*TaskReply) error {
92     c.mu.Lock()
93     defer c.mu.Unlock()
94
95     if !c.mapDone {
96         for i, task := range c.mapTasks {
97             if c.tasks[i].Status == Pending {
98                 c.tasks[i].Status = InProgress
99                 c.tasks[i].WorkerId = args.WorkerId
100                 c.tasks[i].StartTime = time.Now()
101                 reply.Task = task
102                 return nil
103             }
104         }
105         c.mapDone = true
106         for i := 0; i < c.nMap; i++ {
107             if c.tasks[i].Status != Completed {
108                 c.mapDone = false
109                 break
110             }

```

```

111     }
112 }
113
114 if c.mapDone && !c.reduceDone {
115     for i, task := range c.reduceTasks {
116         taskId := c.nMap + i
117         if c.tasks[taskId].Status == Pending {
118             c.tasks[taskId].Status = InProgress
119             c.tasks[taskId].WorkerId = args.WorkerId
120             c.tasks[taskId].StartTime = time.Now()
121             reply.Task = task
122             return nil
123         }
124     }
125     c.reduceDone = true
126     for i := c.nMap; i < c.nMap+c.nReduce; i++ {
127         if c.tasks[i].Status != Completed {
128             c.reduceDone = false
129             break
130         }
131     }
132 }
133
134 reply.Task = Task{Type: None}
135 return nil
136 }
137
138 func (c *Coordinator) TaskDone(args *ReportArgs, reply
139 *ReportReply) error {
140     c.mu.Lock()
141     defer c.mu.Unlock()
142
143     taskId := args.TaskNum
144     if args.TaskType == Reduce {
145         taskId += c.nMap
146     }
147
148     if c.tasks[taskId].WorkerId != args.WorkerId {
149         reply.Success = false
150         return nil
151     }
152
153     c.tasks[taskId].Status = Completed
154     reply.Success = true
155     return nil
156 }

```

2.1.1 Explanation

The `coordinator.go` file is the backbone of the MapReduce system, implementing the `Coordinator` struct and its associated methods. The coordinator is responsible for orchestrating the entire MapReduce workflow, including task assignment, progress tracking, and fault tolerance.

- **Structs:**

- **Task:** Represents a MapReduce task, which could be a map or reduce task. Fields include:
 - **Type:** Indicates whether the task is a map or reduce task.
 - **Num:** The task number.
 - **Filename:** The input file for map tasks (empty for reduce tasks).
 - **NumReduce** and **NumMap:** The number of reduce and map tasks.
 - **TaskArgs:** Contains the worker's ID.
 - **TaskReply:** Contains the task assigned by the coordinator.
 - **ReportArgs:** Contains the worker's ID, task number, and task type (map or reduce).
 - **ReportReply:** Indicates whether the task completion report was successful.

- **Methods:**

- **server():** Initializes an RPC server over a Unix socket (generated by `coordinatorSock()`). It registers the coordinator for RPC calls, removes any existing socket file, and starts an HTTP server in a goroutine to handle worker requests.
- **monitor():** Runs continuously in a goroutine to detect stalled tasks. It locks the mutex, checks each task in progress, and resets tasks to `Pending` if they exceed a 10-second timeout, enhancing fault tolerance by reassigning failed tasks.
- **Done():** Returns `true` when both map and reduce phases are complete, allowing the system to terminate gracefully.
- **MakeCoordinator():** Constructs a new coordinator instance. It initializes map tasks from input files, creates reduce tasks, sets up task trackers, and launches the monitor and server goroutines.
- **AssignTask():** An RPC method that assigns tasks to workers. It prioritizes map tasks until all are completed (`mapDone` becomes `true`), then assigns reduce tasks. If no tasks are available, it returns a `None` task type. The method updates task status and tracks worker assignment.
- **TaskDone():** An RPC method that marks a task as completed. It validates the worker's ID and updates the task status to `Completed`. If the report is valid, it returns `true`, signaling success.

2.2 rpc.go

```
1 package mr
2
3 import (
4     "net/rpc"
5 )
6
7 type Task struct {
8     Type      TaskType
```

```

9      Num      int
10     Filename  string
11     NumReduce int
12     NumMap    int
13 }
14
15 type TaskArgs struct {
16     WorkerId string
17 }
18
19 type TaskReply struct {
20     Task Task
21 }
22
23 type ReportArgs struct {
24     WorkerId string
25     TaskNum  int
26     TaskType TaskType
27 }
28
29 type ReportReply struct {
30     Success bool
31 }
32
33 type TaskType int
34
35 const (
36     Map TaskType = iota
37     Reduce
38     None
39 )
40
41 func call(rpcname string, args interface{}, reply
42         interface{}) bool {
43     client, err := rpc.DialHTTP("unix", coordinatorSock())
44     if err != nil {
45         return false
46     }
47     defer client.Close()
48
49     err = client.Call(rpcname, args, reply)
50     if err != nil {
51         return false
52     }
53     return true
54 }

```

2.2.1 Explanation

The `rpc.go` file defines the data structures used for remote procedure calls (RPCs) between the coordinator and the workers. It also contains the `call()` function to facilitate communication with the coordinator.

- **Structs:**

- **Task:** Represents a MapReduce task, which could be a map or reduce task. Fields include:
- **Type:** Indicates whether the task is a map or reduce task.
- **Num:** The task number.
- **Filename:** The input file for map tasks (empty for reduce tasks).
- **NumReduce** and **NumMap:** The number of reduce and map tasks.
- **TaskArgs:** Contains the worker's ID.
- **TaskReply:** Contains the task assigned by the coordinator.
- **ReportArgs:** Contains the worker's ID, task number, and task type (map or reduce).
- **ReportReply:** Indicates whether the task completion report was successful.
- **TaskType:** An enumeration defining the three possible types of tasks: Map, Reduce, and None.

- **Function:**

- `call()`: This function is used by workers to send RPC requests to the coordinator. It opens a connection to the coordinator via a Unix socket and sends the RPC call. If the call is successful, it returns `true`, otherwise `false`.

2.3 worker.go

```
1 package mr
2
3 import (
4     "fmt"
5     "log"
6     "os"
7     "time"
8 )
9
10 func worker() {
11     for {
12         var args TaskArgs
13         var reply TaskReply
14         if !call("Coordinator.AssignTask", &args, &reply) {
15             log.Fatal("Call to Coordinator failed!")
16         }
17
18         task := reply.Task
19         if task.Type == None {
20             return
21         }
22
23         if task.Type == Map {
```

```

24         // Perform Map Task
25         err := doMapTask(task)
26         if err != nil {
27             log.Printf("Map task failed: %v", err)
28             continue
29         }
30     } else if task.Type == Reduce {
31         // Perform Reduce Task
32         err := doReduceTask(task)
33         if err != nil {
34             log.Printf("Reduce task failed: %v", err)
35             continue
36         }
37     }
38
39     var reportArgs ReportArgs
40     var reportReply ReportReply
41     reportArgs.WorkerId = args.WorkerId
42     reportArgs.TaskNum = task.Num
43     reportArgs.TaskType = task.Type
44
45     if !call("Coordinator.TaskDone", &reportArgs,
46 &reportReply) {
47         log.Fatal("Call to Coordinator failed!")
48     }
49
50     if !reportReply.Success {
51         log.Printf("Failed to report task completion for
52 task %d", task.Num)
53     }
54 }

```

2.3.1 Explanation

The `worker.go` file implements the worker's behavior. It continuously requests tasks from the coordinator, processes them, and reports the completion back to the coordinator.

- `worker()`: The worker function enters an infinite loop where it:
 - Calls the `Coordinator.AssignTask` RPC to request a task.
 - If the task is `None`, it terminates.
 - Depending on the task type (map or reduce), it calls the respective function (`doMapTask()` or `doReduceTask()`) to perform the task.
 - After completing the task, it calls `Coordinator.TaskDone` to report the completion of the task.

3 Test Output and Analysis

The output generated by running the test script `Test` demonstrates the proper functioning of the MapReduce system, as each worker completes its assigned task (either map or reduce) and reports the completion to the coordinator.

```
saeid@saeid:~/Desktop/Distributed/1/Project 1/Project 1/Distributed_Systems_CA1$ cd src/main
saeid@saeid:~/Desktop/Distributed/1/Project 1/Project 1/Distributed_Systems_CA1/src/main$ bash test-mr.sh
*** Starting wc test.
--- wc test: PASS
2025/04/18 03:27:44 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:27:44 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:27:44 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
*** Starting indexer test.
2025/04/18 03:27:48 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:27:48 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- indexer test: PASS
*** Starting map parallelism test.
2025/04/18 03:27:55 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:27:55 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- map parallelism test: PASS
*** Starting reduce parallelism test.
2025/04/18 03:28:03 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:28:03 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- reduce parallelism test: PASS
*** Starting job count test.
2025/04/18 03:28:21 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:28:21 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:28:21 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- job count test: PASS
2025/04/18 03:28:22 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
*** Starting early exit test.
2025/04/18 03:28:29 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:28:29 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:28:29 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- early exit test: PASS
*** Starting crash test.
2025/04/18 03:29:11 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:29:12 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
2025/04/18 03:29:12 dialing:dial unix /var/tmp/5840-mr-1000: connect: connection refused
--- crash test: PASS
*** PASSED ALL TESTS
saeid@saeid:~/Desktop/Distributed/1/Project 1/Project 1/Distributed_Systems_CA1/src/main$
```

Figure 1: Test Output